

# Recitation 2

## Gradient Descent and Stochastic Gradient Descent

Yanlai Yang

CDS

February 1, 2023

# Agenda

- Gradient Descent
  - Adaptive Learning Rate
- Stochastic Gradient Descent
- Application
  - Linear Regression
  - Logistic Regression

# Gradient Descent Recap

## Gradient Descent

- Initialize  $x = 0$
  - Repeat:
    - $x \leftarrow x - \underbrace{\eta}_{\text{step size}} \nabla f(x)$
  - Until stopping criterion satisfied
- 
- Choosing the step size is the key in gradient descent
  - A fixed step size will work, eventually, as long as it's small enough


# Gradient Descent

## Gradient Descent Algorithm

- Goal: find  $\theta^* = \arg \min_{\theta} J(\theta)$
- $\theta^0 :=$  [initial condition]
- $i := 0$
- while not [termination condition]:
  - compute  $\nabla J(\theta_i)$
  - $\epsilon_i :=$  [choose learning rate at iteration  $i$ ]
  - $\theta^{i+1} := \theta^i - \epsilon_i \nabla J(\theta_i)$
  - $i = i + 1$
- return  $\theta^i$

# Gradient Descent

- How to initialize  $\theta^0$ ?
  - sample from some distribution
  - compose  $\theta^0$  using some heuristics
- How to choose termination conditions?
  - run for a fixed number of iteration
  - the value of  $f(\theta)$  stabilizes
  - $\theta^i$  converges
- What is a good learning rate?





# Learning Rate

## Application

Suppose we would like to find  $\theta^* \in \mathbb{R}$  that minimizes  $f(\theta) = \theta^2 - 2\theta + 1$ . The gradient (in this case, the derivative)  $\nabla f(\theta) = 2\theta - 2$ . We can easily see that  $\theta^* = \operatorname{argmin}_{\theta} f(\theta) = 1$ .

# Learning Rate

- We applied gradient descent for 1000 iterations on  $f(\theta) = \theta^2 - 2\theta + 1$  with varying learning rate  $\epsilon \in \{1, 0.1, 0.01, 0.001, 0.0001\}$
- When the learning rate is too large ( $\epsilon = 1$ ),  $f(\theta)$  does not decrease through iterations. The value of  $\theta_i$  at each iteration significantly fluctuates.
- When the learning rate is too small ( $\epsilon = 0.0001$ ),  $f(\theta)$  decreases very slowly.



# Adaptive Learning Rate

- Instead of using a fixed learning rate through all iterations, we can adjust our learning rate in each iteration using a simple algorithm.
- At each iteration  $i$ :
  - $\tilde{\theta} := \theta_{i-1} - \epsilon_{i-1} \nabla f(\theta_{i-1})$
  - $\delta := f(\theta_{i-1}) - f(\tilde{\theta})$
  - if  $\delta \geq \text{threshold}$ :
    - we achieve a satisfactory reduction on  $f(\theta)$
    - $\theta_i = \tilde{\theta}$
    - maybe we can consider increasing the learning rate for next iteration  
 $\epsilon_j := 2\epsilon_{i-1}$
  - else:
    - the reduction is unsatisfactory
    - $\theta_i = \theta_{i-1}$
    - the learning rate is too large, so we reduce the learning rate
    - $\epsilon_i := \frac{1}{2}\epsilon_{i-1}$



# Adaptive Learning Rate

How to decide a proper threshold for  $f(\theta_{i-1}) - f(\tilde{\theta})$

## Armijo rule

If learning rate  $\epsilon$  satisfies

$$f(\theta_{i-1}) - f(\tilde{\theta}) \geq \frac{1}{2}\epsilon \|\nabla f(\theta_{i-1})\|^2$$

then  $f(\theta)$  is guaranteed to converge to a (local) minimum under certain technical assumptions.

You can find more details at **this link**

# Stochastic Gradient Descent

## Stochastic Gradient Descent

- Initialize  $w = 0$
  - Repeat:
    - randomly choose training point  $(x_i, y_i) \in \mathcal{D}_n$
    - $w \leftarrow w - \eta \underbrace{\nabla_w \ell(f_w(x_i), y_i)}_{\text{Grad(Loss on i'th example)}}$
  - Until stopping criterion satisfied
- 
- Equivalent to Minibatch Gradient Descent with batch size  $N = 1$ .
  - Use a single randomly chosen point to determine step direction.

# Minibatch Gradient Descent

## Minibatch Gradient Descent (minibatch size $N$ )

- Initialize  $w = 0$
  - Repeat:
    - randomly choose  $N$  points  $\{(x_i, y_i)\}_{i=1}^N \subset \mathcal{D}_n$
    - $w \leftarrow w - \eta \left[ \frac{1}{N} \sum_{i=1}^N \nabla_w \ell(f_w(x_i), y_i) \right]$
  - Until stopping criterion satisfied
- 
- Minibatch gradient is an unbiased estimate of full-batch gradient:
$$\mathbb{E} \left[ \nabla \hat{R}_N(w) \right] = \nabla \hat{R}_n(w)$$
  - Use a random subset of size  $N$  to determine step direction
    - Bigger  $N$ : Better estimate of the gradient, but slower (more data to touch)
    - Smaller  $N$ : Worse estimate of the gradient, but faster

# Gradient Descent for Linear Regression

## Linear Least Squares Regression Setup

- **Data:** Inputs are feature vectors of dimension  $d$ . Outputs are continuous scalars.

$$\mathcal{D} = \left\{ \mathbf{x}^{(i)}, y^{(i)} \right\}_{i=1}^n \text{ where } \mathbf{x} \in \mathbb{R}^d \text{ and } y \in \mathbb{R}$$

- **Hypothesis Space:**  $\mathcal{F} = \{ f : \mathbb{R}^d \rightarrow \mathbb{R} \mid f(\mathbf{x}) = \theta^T \mathbf{x}, \theta \in \mathbb{R}^d \}$
- **Action:** Our prediction is a linear function of the inputs

$$\begin{aligned} \hat{y} &= f_{\theta}(\mathbf{x}) = \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_d x_d \\ \hat{y} &= f_{\theta}(\mathbf{x}) = \theta^T \mathbf{x} \end{aligned} \quad (\text{We assume } x_1 \text{ is } 1)$$

- **Loss:**  $\ell(\hat{y}, y) = (y - \hat{y})^2$

# Gradient Descent for Linear Regression

- **Goal:** Finding the set of parameters that minimize the empirical risk:

$$\hat{R}_n(\theta) = \frac{1}{n} \sum_{i=1}^n \left( \theta^T \mathbf{x}^{(i)} - y^{(i)} \right)^2$$

where  $\theta \in R^d$  parameterizes the hypothesis space  $\mathcal{F}$

- Set our cost function:

$$J(\theta) = \frac{1}{2} \sum_{i=1}^n \left( \theta^T \mathbf{x}^{(i)} - y^{(i)} \right)^2$$

# Three Approach to solving $\theta^* = \underset{\theta}{\operatorname{argmin}} J(\theta)$

- **Approach 1:** Closed Form Solution (set derivatives equal to zero and solve for parameters)
  - pros: one shot algorithm!
  - cons: does not scale to large datasets (matrix inverse is bottleneck)
- **Approach 2:** Gradient Descent (take larger, more certain steps toward the negative gradient)
  - pros: conceptually simple, guaranteed convergence
  - cons: batch, often slow to converge
- **Approach 3:** Stochastic Gradient Descent (take many small, quick steps opposite the gradient)
  - pros: memory efficient, fast convergence, less prone to local optima
  - cons: convergence in practice requires tuning and fancier variants

# Approach 1: Close-Form Solution

- Transform the cost function in matrix form

$$\begin{aligned} J(\theta) &= \frac{1}{2} \sum_{i=1}^n \left( \mathbf{x}_i^T \theta - y_i \right)^2 \\ &= \frac{1}{2} (X\theta - \bar{y})^T (X\theta - \bar{y}) = \frac{1}{2} \|X\theta - y\|_2^2 \end{aligned}$$

- To minimize  $J(\theta)$ , take derivative and set to zero:

$$\begin{aligned} \nabla J(\theta) &= (X^T X \theta - X^T y) = X^T (X\theta - y) = 0 \\ \Rightarrow \hat{\theta} &= (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y} \end{aligned}$$

- Ensure invertibility of  $\mathbf{X}^T \mathbf{X}$
- What if  $X$  has less than full column rank?

## Approach 2: Iterative Method GD

### Gradient Descent Algorithm

- $\theta^0 :=$  [initial condition]
- $i := 0$
- while not [termination condition]:
  - $\theta^{i+1} := \theta^i - \epsilon_i \nabla_{\theta} J(\theta)$
  - $i = i + 1$
- return  $\theta^i$

Recall:

$$\nabla_{\theta} J(\theta) = \begin{bmatrix} \frac{d}{d\theta_1} J(\theta) \\ \frac{d}{d\theta_2} J(\theta) \\ \vdots \\ \frac{d}{d\theta_d} J(\theta) \end{bmatrix}$$




# Approach 3: Iterative Method SGD


## Stochastic Gradient Descent Algorithm

- $\theta^0 :=$  [initial condition]
- $i := 0$
- while not [termination condition]:
  - For each training pair  $(x^j, y^j)$  (in random order)
  - $\theta^{i+1} := \theta^i - \epsilon_i \nabla_{\theta} J^{(j)}(\theta)$  {with  $J^{(j)}(\theta) = \frac{1}{2} \left( \theta^T \mathbf{x}^{(j)} - y^{(j)} \right)^2$ }
  - $i = i + 1$
- return  $\theta^i$

GD



SGD



# Gradient Descent for Logistic Regression

## Binary Classification Setup for Logistic Regression

- **Data:** Inputs are feature vectors of dimension  $d$ . Targets are class labels.  $\mathcal{D} = \{\mathbf{x}^{(i)}, y^{(i)}\}_{i=1}^n$  where  $\mathbf{x} \in \mathbb{R}^d$  and  $y \in \{0, 1\}$
- **Action:** Our prediction is the probability of class label given linear signals

$$h_{\theta}(x) = g(\theta^T x) = \frac{1}{1 + e^{-\theta^T x}} \quad \text{with} \quad g(z) = \frac{1}{1 + e^{-z}}$$

- Sigmoid Function  $g(z)$ : takes a real-valued number and maps it into the range  $[0, 1]$  (Probability Interpretation)
- Assume

$$\begin{cases} P(y = 1 \mid x; \theta) = h_{\theta}(x) \\ P(y = 0 \mid x; \theta) = 1 - h_{\theta}(x) \end{cases}$$

- More Compactly:  $p(y \mid x; \theta) = (h_{\theta}(x))^y (1 - h_{\theta}(x))^{1-y}$

# Learning Logistic Regression

- Assumption:  $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)$  are independently generated
- Likelihood:** The probability of getting the  $y_1, \dots, y_n$  in  $\mathcal{D}$  from the corresponding  $\mathbf{x}_1, \dots, \mathbf{x}_n$

$$\begin{aligned}
 P(y_1, \dots, y_n \mid \mathbf{x}_1, \dots, \mathbf{x}_n) &= \prod_{i=1}^n p(y^{(i)} \mid x^{(i)}; \theta) \\
 &= \prod_{i=1}^n (h_{\theta}(x^{(i)}))^{y^{(i)}} (1 - h_{\theta}(x^{(i)}))^{1-y^{(i)}}
 \end{aligned}$$

- Goal:** maximize the log likelihood (Easier)

$$\ell(\theta) = \log L(\theta) = \sum_{i=1}^n y^{(i)} \log h_{\theta}(x^{(i)}) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)}))$$

- Equivalent to minimize the objective function with logistic loss:

$$J(\theta) = \sum_{i=1}^n \ell(h_{\theta}(x^{(i)}), y^{(i)})$$

$$\text{where } \ell(h(x), y) = -y \log(h_{\theta}(x)) - (1 - y) \log(1 - h_{\theta}(x))$$

# Learning Logistic Regression

- Analytic solution won't work
- Find optimum using iterative methods: Gradient Ascent or Stochastic Gradient Ascent
  - Gradient ascent rule:  $\theta := \theta + \alpha \nabla_{\theta} \ell(\theta)$
  - Stochastic gradient ascent rule:  $\theta := \theta + \alpha \nabla_{\theta} \ell^{(i)}(\theta)$  for random training pair  $(x^i, y^i)$
- For one training example  $(x, y)$ , the partial derivative of log likelihood  $\ell(\theta)$ :

$$\begin{aligned}
 &= \left( y \frac{1}{g(\theta^T x)} - (1 - y) \frac{1}{1 - g(\theta^T x)} \right) \frac{\partial}{\partial \theta_j} g(\theta^T x) \\
 &= \left( y \frac{1}{g(\theta^T x)} - (1 - y) \frac{1}{1 - g(\theta^T x)} \right) g(\theta^T x) (1 - g(\theta^T x)) \frac{\partial}{\partial \theta_j} \theta^T x \\
 &= (y(1 - g(\theta^T x)) - (1 - y)g(\theta^T x)) x_j \\
 &= (y - h_{\theta}(x)) \cdot x_j
 \end{aligned}$$

# References

- DS-GA 1003 Machine Learning Spring 2021 & 2022
- Stanford CS229 Note1
- CMU 10-701 Linear Regression Slide