

Neural Network and Backpropagation Questions

Daeyoung Kim, Xintian Han

CDS, NYU

April 21, 2021

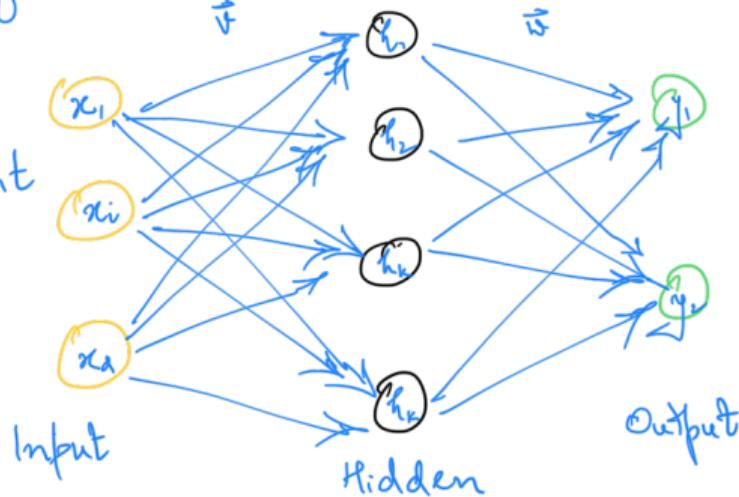
[TK] Annotated

Recap: Neural Network and Backpropagation

- + Feature / representation learning
- + Gradient Descent
- + Special case of automatic differentiation
- + Algorithm to compute gradient

- Initial idea from neuroscience

Single hidden layer is sufficient to approximate any function



partial derivatives + chain rule

$$h_k = \sigma(t_k^T x)$$

$$f(x) = \sum_{k=1}^K w_k h_k(x)$$

$$= \sum_{k=1}^K w_k \sigma(t_k^T x)$$

Question 1: Step Activation Function ¹

Suppose we have a neural network with one hidden layer.

$$f(x) = w_0 + \sum_i w_i h_i(x); \quad h_i(x) = g(b_i + v_i x),$$

where activation function g is defined as

$$\underline{g(z)} = \begin{cases} 1 & \text{if } z \geq 0 \\ 0 & \text{if } z < 0 \end{cases}$$

step fn

Which of the following functions can be exactly represented by this neural network?

- polynomials of degree one: $I(x) = ax + b$
- hinge loss: $I(x) = \max(1 - x, 0)$
- polynomials of degree two: $I(x) = ax^2 + bx + c$
- piecewise constant functions

¹From CMU

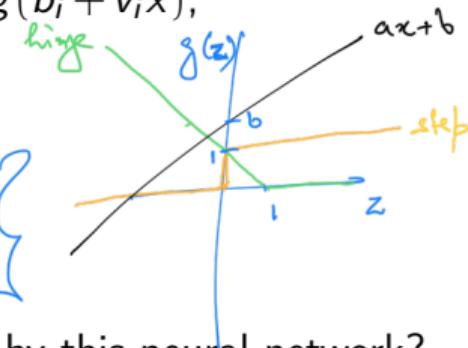
[Solution] Question 1: Step Activation Function

Suppose we have a neural network with one hidden layer.

$$f(x) = w_0 + \sum_i w_i h_i(x); \quad h_i(x) = g(b_i + v_i x),$$

where activation function g is defined as

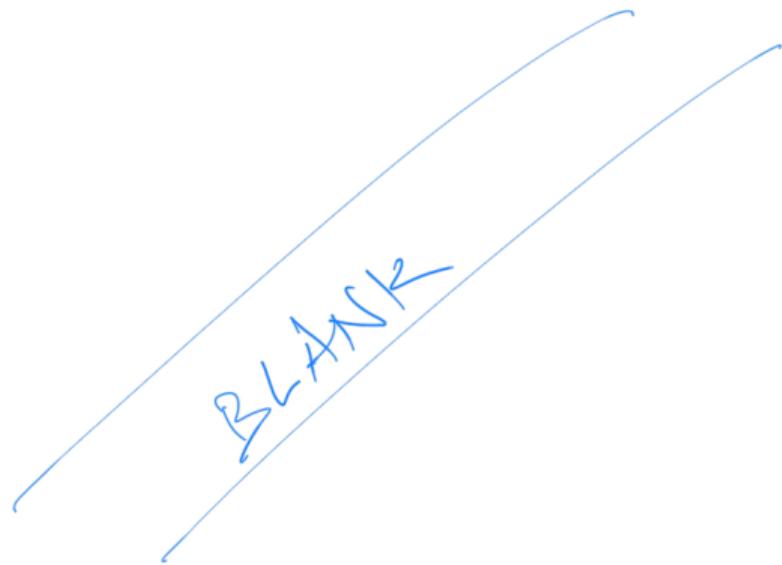
$$g(z) = \begin{cases} 1 & \text{if } z \geq 0 \\ 0 & \text{if } z < 0 \end{cases}$$



Which of the following functions can be exactly represented by this neural network?

- polynomials of degree one: $I(x) = ax + b$ **No** — $f(x) = w_0 + \sum_i w_i \mathbb{1}(h_i(x) \geq 0)$
If g can be identity function, then the answer is **Yes** — $f(x) = w_0 + wv x + wb$
 $\Rightarrow a = wv, b = w_0 + wb$
- hinge loss: $I(x) = \max(1 - x, 0)$ **No**
- polynomials of degree two: $I(x) = ax^2 + bx + c$ **No**
- piecewise constant functions **Yes**
 $(-c) \cdot g(x - b) + (c) \cdot g(x - a)$ can represent $I(x) = c, a \leq x < b$.

[Solution] Question 1: Step Activation Function



Question 2: Power of ReLU²

Rectified Linear Unit

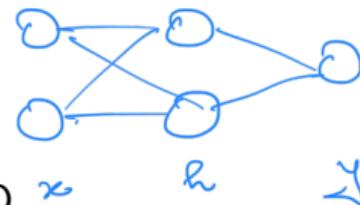
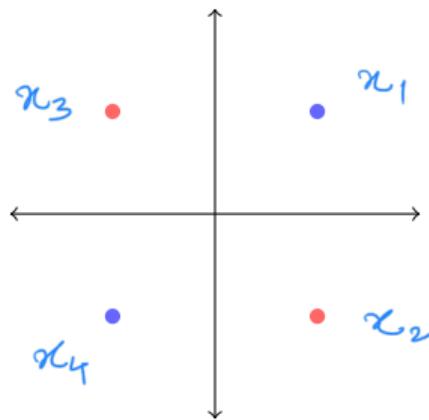
Consider the following small NN:

$$w_2^\top \text{ReLU}(W_1 x + b_1) + b_2$$

where the data is 2D, W_1 is 2 by 2, b_1 is 2D, w_2 is 2D and b_2 is 1D.

$$x_1 = (1, 1) \quad y_1 = 1; \quad x_2 = (1, -1) \quad y_2 = -1; \quad x_3 = (-1, 1) \quad y_3 = -1; \quad x_4 = (-1, -1) \quad y_4 = 1$$

Find b_1, b_2, W_1, w_2 to solve the problem. (Separate points from class $y = 1$ and $y = -1$.)

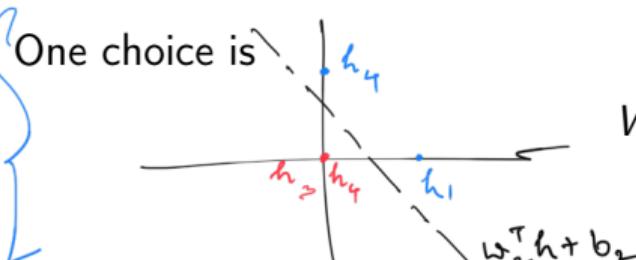
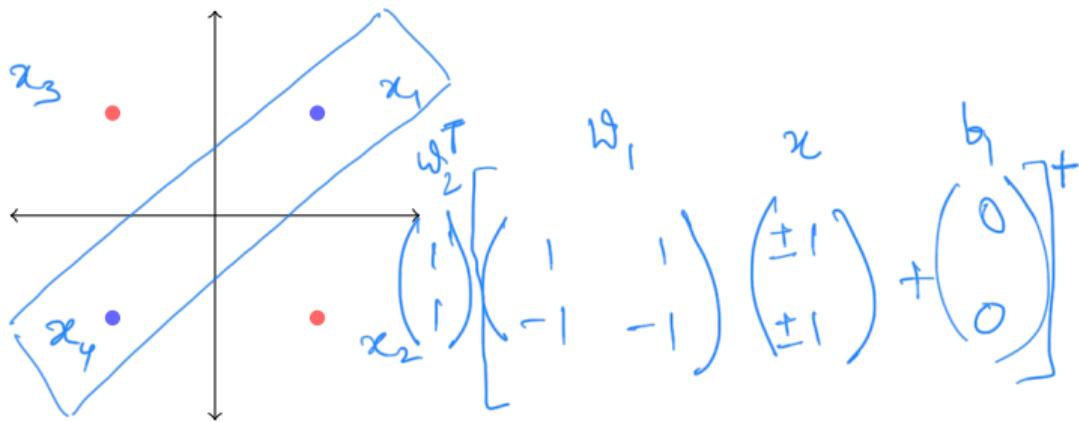
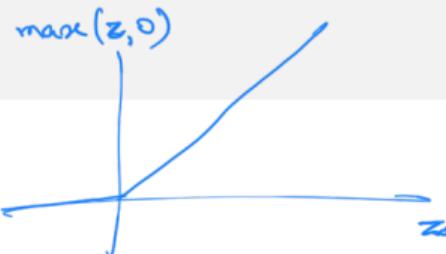


²From Harvard

[Solution] Question 2: Power of ReLU

$$h \quad f(x) \quad \left\{ \begin{array}{l} z \\ 0 \end{array} \right. \quad \begin{array}{l} z > 0 \\ z \leq 0 \end{array}$$

x_1	1	1	$[2 \ 0]$	$ w_2^T \text{ReLU}(W_1x + b_1) + b_2 $
x_2	1	-1	$[0 \ 0]$	-1
x_3	-1	1	$[0 \ 0]$	-1
x_4	-1	-1	$[0 \ 2]$	1



$$W_1 = \begin{pmatrix} 1 & 1 \\ -1 & -1 \end{pmatrix}, b_1 = \begin{pmatrix} 0 \\ 0 \end{pmatrix}$$

$$w_2 = \begin{pmatrix} 1 \\ 1 \end{pmatrix}, b_2 = -1$$

$$+ (-1) \\ b_2$$

[Solution] Question 2: Power of ReLU

Idea: Map points into different space after performing various operations so as to make them separable / be able to classify

e.g. linear transformations / matrix multiplications

Combination
 $W = U\Sigma V^T$
Singular Value Decomposition

Rotation: orthonormal matrix
Scaling: diagonal matrix
Reflection: matrix determinant -ve

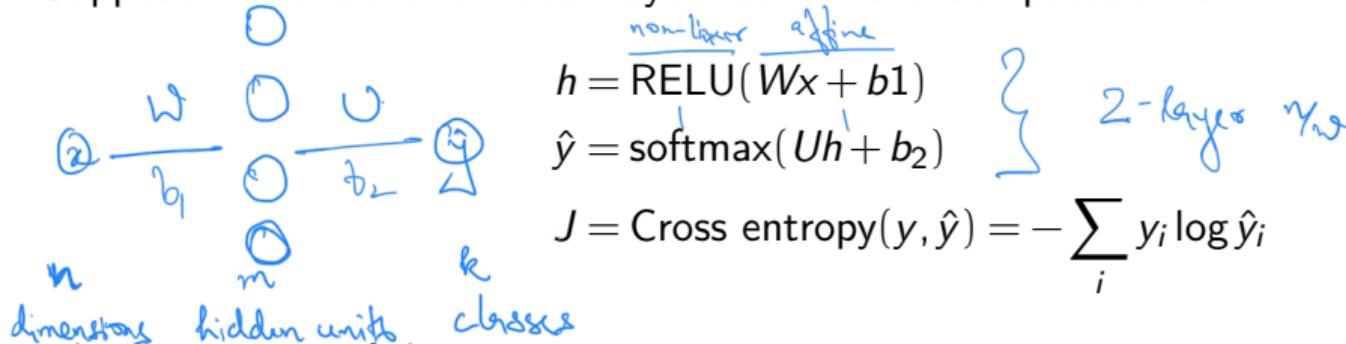
Affine - translation , non-linear transformations
(+b)

Shearing
Projection
Squeeze

Reference : <https://atcold.github.io/pytorch-Deep-Learning/en/week01/01-3/>

Question 3: Backpropagation ³

Suppose we have a one hidden layer network and computation is:



The dimensions of the matrices are:

$$W \in \mathbb{R}^{m \times n} \quad x \in \mathbb{R}^n \quad b_1 \in \mathbb{R}^m \quad U \in \mathbb{R}^{k \times m} \quad b_2 \in \mathbb{R}^k$$

Use backpropagation to calculate these four gradients

$$\frac{\partial J}{\partial b_2} \quad \frac{\partial J}{\partial U} \quad \frac{\partial J}{\partial b_1} \quad \frac{\partial J}{\partial W}$$

³From Stanford

[Solution] Question 3: Backpropagation

$$\hat{y}_i = \text{softmax}(z_2^{(i)}) = \frac{e^{z_2^{(i)}}}{\sum e^{z_2^{(i)}}}$$

$$\hat{y} = \text{softmax}(z_2)$$

$$\frac{\partial J}{\partial z_2} \cdot \frac{\partial z_2}{\partial b_2} \rightarrow \frac{\partial J}{\partial b_2} = \delta_1$$

$$\frac{\partial J}{\partial y} \cdot \frac{\partial y}{\partial z_2} \cdot \frac{\partial z_2}{\partial U} \rightarrow \frac{\partial J}{\partial U} = \delta_1 h^T$$

$$h = \text{ReLU}(z_1)$$

$$\frac{\partial J}{\partial z_1} \cdot \frac{\partial z_1}{\partial b_1} \rightarrow \frac{\partial J}{\partial b_1} = \delta_2$$

$$\frac{\partial J}{\partial W} = \delta_2 x^T$$

$$z_2 = Uh + b_2 \quad \delta_1 = \frac{\partial J}{\partial z_2} = \hat{y} - y$$

$$\delta_1 = \frac{\partial J}{\partial y} \cdot \frac{\partial y}{\partial z_2}$$

↓

Cross-entropy \Rightarrow softmax

$$\frac{\partial J}{\partial h} \frac{\partial h}{\partial z_1}$$

$$\frac{\partial J}{\partial z_2} \cdot \frac{\partial y}{\partial z_2} \cdot \frac{\partial h}{\partial h} \cdot \frac{\partial h}{\partial z_1}$$

δ_1 U^T

[Solution] Question 3: Backpropagation

Software

$$\frac{\partial \hat{y}_i}{\partial z_2^{(i)}} = \frac{\partial}{\partial z_2^{(i)}} \left[\frac{e^{z_2^{(i)}}}{\sum_j e^{z_2^{(j)}}} \right]$$

Quotient rule

$$\text{for } f(x) = \frac{g(x)}{h(x)}$$

$$\rightarrow f'(x) = \frac{g'(x)h(x) - h'(x)g(x)}{[h(x)]^2}$$

$i \neq j$

$$= \frac{e^{z_2^{(i)}} \sum - e^{z_2^{(i)}} e^{z_2^{(j)}}}{\sum^2}$$

$$= \frac{e^{z_2^{(i)}}}{\sum} - \left(\frac{e^{z_2^{(i)}}}{\sum} \right)^2$$

$$\frac{\partial \hat{y}_i}{\partial z_2^{(i)}} = \hat{y}_i - \hat{y}_i \hat{y}_i = \hat{y}_i (1 - \hat{y}_i)$$

$$= \frac{0 \sum - e^{z_2^{(i)}} e^{z_2^{(i)}}}{\sum^2} = - \frac{e^{z_2^{(i)}} e^{z_2^{(i)}}}{\sum^2} = - \hat{y}_i \hat{y}_i$$

Intuition

Coding Exercise

Cross-Entropy

$$\frac{\partial J}{\partial \hat{y}_i} = \frac{\partial}{\partial \hat{y}_i} \left[- \sum_i y_i \log \hat{y}_i \right]$$

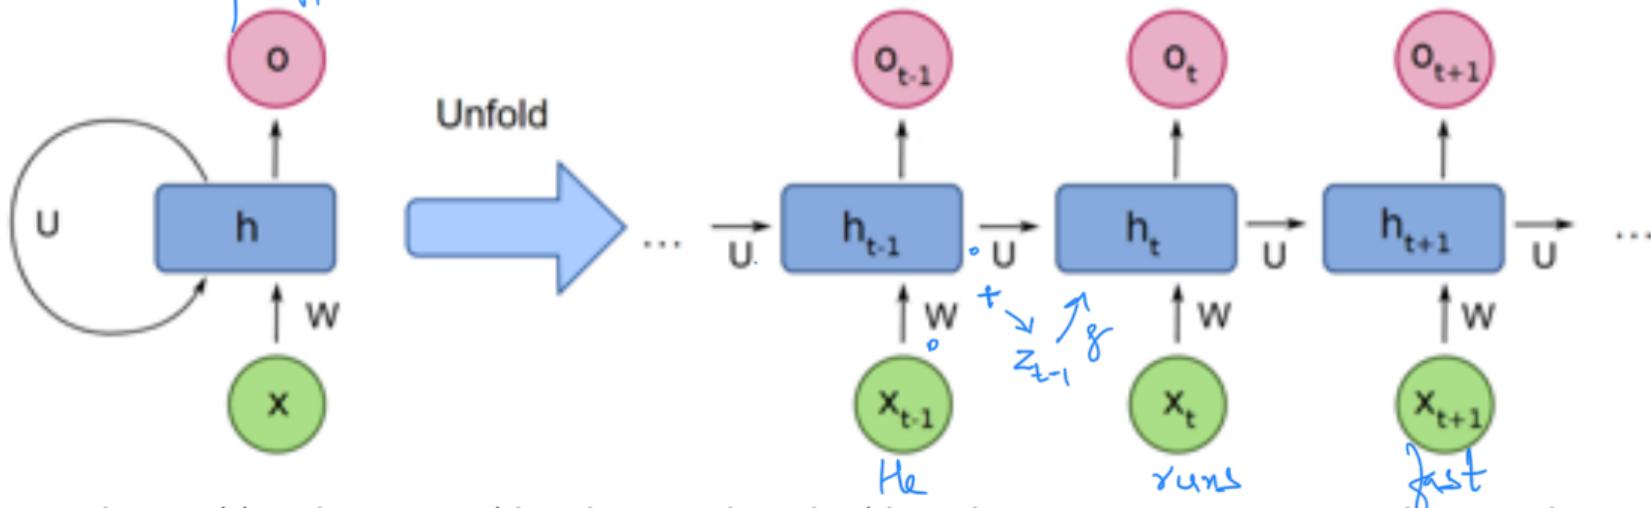
$$= - \frac{\hat{y}_i}{y_i}$$

- Computation graph hands-on

$$\begin{aligned}\frac{\partial J}{\partial z_i^{(l)}} &= \frac{\partial J}{\partial \hat{y}_i} \cdot \frac{\partial \hat{y}_i}{\partial z_i^{(l)}} = - \frac{\hat{y}_i}{y_i} \frac{y_i (1 - \hat{y}_i)}{\hat{y}_i} - \sum_{j \neq i} \frac{y_j}{\hat{y}_j} (-\hat{y}_j) \\ &= -\hat{y}_i + \hat{y}_i \hat{y}_i - \sum_{j \neq i} y_j (-\hat{y}_j) \\ &= \frac{\hat{y}_i (y_i + \sum_{j \neq i} y_j)}{\sum y_j = 1} - \hat{y}_i = \hat{y}_i - y_i\end{aligned}$$

[Optional] Recurrent Neural Networks

- + For processing variable length i/p, e.g. text-sentences of diff len
- + Applies recursive fn to each element, e.g. words of sentence



Source: <https://medium.com/deeplearningbrasilia/deep-learning-recurrent-neural-networks-f9482a24d010>

- + Also helps carry over information from prior time steps/position through representations learnt

[Optional]: Backpropagation in RNN

+ BLTT: through time

Suppose we have a recurrent neural network (RNN). The recursive function is:

$$\begin{aligned}\mathbf{z}_{t-1} &= \mathbf{Wx}_{t-1} + \mathbf{Uh}_{t-1}, \\ \mathbf{h}_t &= g(\mathbf{z}_{t-1}),\end{aligned}$$

where \mathbf{h}_t is the hidden state and \mathbf{x}_t is the input at time step t . \mathbf{W} and \mathbf{U} are the weighted matrix. g is an element-wise activation function. And \mathbf{h}_0 is a given fixed initial hidden state.

- Assume loss function \mathcal{L} is a function of \mathbf{h}_T . Given $\partial \mathcal{L} / \partial \mathbf{h}_T$, calculate $\partial \mathcal{L} / \partial \mathbf{U}$ and $\partial \mathcal{L} / \partial \mathbf{W}$.
- Suppose g' is always greater than λ and the smallest singular value of \mathbf{U} is larger than $1/\lambda$. What will happen to the gradient $\partial \mathcal{L} / \partial \mathbf{U}$ and $\partial \mathcal{L} / \partial \mathbf{W}$?
- Suppose g' is always smaller than λ and the largest singular value of \mathbf{U} is smaller than $1/\lambda$. What will happen to the gradient $\partial \mathcal{L} / \partial \mathbf{U}$ and $\partial \mathcal{L} / \partial \mathbf{W}$?

[Solution] [Optional]: Backpropagation in RNN

$$\begin{aligned}\frac{\partial L}{\partial U} &= \sum_t \frac{\partial L}{\partial U} ; \quad \frac{\partial L^{[T]}}{\partial U} = \frac{\partial L}{\partial h_T} \cdot \frac{\partial h_T}{\partial z_{T-1}} \cdot \frac{\partial}{\partial U} (W_{z_{T-1}} + U_{h_{T-1}}) \\ &= \frac{\partial L}{\partial h_T} \cdot \frac{\partial h_T}{\partial z_{T-1}} h_{T-1}^\top + \frac{\partial L}{\partial h_T} \cdot \frac{\partial h_T}{\partial z_{T-1}} U^\top \cdot \underbrace{\left[\frac{\partial h_{T-1}}{\partial U} \right]}_{g'(z_{T-1})} + \dots\end{aligned}$$

[Solution] [Optional]: Backpropagation in RNN

-

$$\frac{\partial \mathcal{L}}{\partial U} = \sum_{t=1}^T (\Pi_{k=t-1}^{T-1} (\mathbf{U}^T D_k)) \frac{\partial \mathcal{L}}{\partial \mathbf{h}_T} \mathbf{h}_{t-1}^T$$

$$\frac{\partial \mathcal{L}}{\partial W} = \sum_{t=1}^T (\Pi_{k=t-1}^{T-1} (\mathbf{U}^T D_k)) \frac{\partial \mathcal{L}}{\partial \mathbf{h}_T} \mathbf{x}_{t-1}^T$$

$D_k = \text{diag}(g'(\mathbf{z}_k))$ is the Jacobian matrix of the element-wise activation function.

- The smallest singular value of the $\mathbf{U}^T D_{k-1}$ will be greater than one. So the smallest singular value of the gradient $\frac{\partial h_s}{\partial h_t}$ will be larger than a^{s-t} for some $a > 1$. So the gradient is going to be exponentially large. This is called exploding gradient.
- The largest singular value of the $\mathbf{U}^T D_{k-1}$ will be smaller than one. So the largest singular value of the gradient $\frac{\partial h_s}{\partial h_t}$ will be smaller than a^{s-t} for some $a < 1$. So the gradient is going to be exponentially small. This is called vanishing gradient.