

# Adaboost

He He

Slides based on Lecture 11c from David Rosenberg's course materials  
(<https://github.com/davidrosenberg/mlcourse>)

CDS, NYU

April 6, 2021

# Boosting

**Bagging** Reduce variance of a low bias, high variance estimator by ensembling many estimators trained in parallel.

**Boosting** Reduce the error rate of a high bias estimator by ensembling many estimators trained in sequential.

- A **weak/base learner** is a classifier that does slightly better than chance.
- Weak learners are like “rules of thumb”:
  - “Viagra”  $\implies$  spam
  - From a friend  $\implies$  not spam
- **Key idea:**
  - Each weak learner focuses on different examples (*reweighted data*)
  - Weak learners have different contributions to the final prediction (*reweighted classifier*)

# AdaBoost: Setting

- *Binary* classification:  $\mathcal{Y} = \{-1, 1\}$
- Base hypothesis space  $\mathcal{H} = \{h : \mathcal{X} \rightarrow \{-1, 1\}\}$ .
- Typical base hypothesis spaces:
  - **Decision stumps** (tree with a single split)
  - Trees with few terminal nodes
  - Linear decision functions

# Weighted Training Set

Each base learner is trained on weighted data.

- Training set  $\mathcal{D} = ((x_1, y_1), \dots, (x_n, y_n))$ .
- Weights  $(w_1, \dots, w_n)$  associated with each example.
- **Weighted empirical risk:**


$$\hat{R}_n^w(f) \stackrel{\text{def}}{=} \frac{1}{W} \sum_{i=1}^n w_i \ell(f(x_i), y_i) \quad \text{where } W = \sum_{i=1}^n w_i$$

- Examples with larger weights have more influence on the loss.

# AdaBoost - Rough Sketch

- Training set  $\mathcal{D} = ((x_1, y_1), \dots, (x_n, y_n))$ .
- Start with equal weight on all training points  $w_1 = \dots = w_n = 1$ .
- Repeat for  $m = 1, \dots, M$ :
  - Find base classifier  $G_m(x)$  that tries to fit weighted training data (but may not do that well)
  - Increase weight on the points  $G_m(x)$  misclassifies
- So far, we've generated  $M$  classifiers:  $G_1, \dots, G_M : \mathcal{X} \rightarrow \{-1, 1\}$ .

# AdaBoost: Schematic



From ESL Figure 10.1

# AdaBoost - Rough Sketch

- Training set  $\mathcal{D} = \{(x_1, y_1), \dots, (x_n, y_n)\}$ .
- Start with equal weight on all training points  $w_1 = \dots = w_n = 1$ .
- Repeat for  $m = 1, \dots, M$ :
  - Base learner fits weighted training data and returns  $G_m(x)$
  - Increase weight on the points  $G_m(x)$  misclassifies
- Final prediction  $G(x) = \text{sign} \left[ \sum_{m=1}^M \alpha_m G_m(x) \right]$ . (recall  $G_m(x) \in \{-1, 1\}$ )
- What are desirable  $\alpha_m$ 's?
  - nonnegative
  - larger when  $G_m$  fits its weighted  $\mathcal{D}$  well
  - smaller when  $G_m$  fits weighted  $\mathcal{D}$  less well



# Adaboost: Weighted Classification Error

- Weights of base learners depend on their performance. How to evaluate each base learner?
- In round  $m$ , base learner gets a weighted training set.
  - Returns a base classifier  $G_m(x)$  that minimizes weighted 0–1 error.
- The **weighted 0-1 error** of  $G_m(x)$  is

$$\text{err}_m = \frac{1}{W} \sum_{i=1}^n w_i 1(y_i \neq G_m(x_i)) \quad \text{where } W = \sum_{i=1}^n w_i.$$

- Notice:  $\text{err}_m \in [0, 1]$ .

# AdaBoost: Classifier Weights

- The weight of classifier  $G_m(x)$  is  $\alpha_m = \ln \left( \frac{1 - \text{err}_m}{\text{err}_m} \right)$ .



- Higher weighted error  $\implies$  lower weight
- When is  $\alpha_m < 0$ ?

## Adaboost: Example Reweighting

- We train  $G_m$  to minimize weighted error, and it achieves  $err_m$ .
- Then  $\alpha_m = \ln\left(\frac{1-err_m}{err_m}\right)$  is the weight of  $G_m$  in final ensemble.

We want the base learner to focus more on examples misclassified by the previous learner.

- Suppose  $w_i$  is weight of example  $i$  before training:
  - If  $G_m$  classifies  $x_i$  correctly, then  $w_i$  is unchanged.
  - Otherwise,  $w_i$  is increased as

$$\begin{aligned}w_i &\leftarrow w_i e^{\alpha_m} \\ &= w_i \left( \frac{1 - err_m}{err_m} \right)\end{aligned}$$

- For  $err_m < 0.5$  (weak learner), this always increases the weight.

# AdaBoost: Algorithm

Given training set  $\mathcal{D} = \{(x_1, y_1), \dots, (x_n, y_n)\}$ .

- ① Initialize observation weights  $w_i = 1, i = 1, 2, \dots, n$ .
- ② For  $m = 1$  to  $M$ :
  - ① Base learner fits weighted training data and returns  $G_m(x)$
  - ② Compute *weighted empirical 0-1 risk*:

$$\text{err}_m = \frac{1}{W} \sum_{i=1}^n w_i \mathbf{1}(y_i \neq G_m(x_i)) \quad \text{where } W = \sum_{i=1}^n w_i.$$

- ③ Compute *classifier weight*:  $\alpha_m = \ln \left( \frac{1 - \text{err}_m}{\text{err}_m} \right)$ .
  - ④ Update *example weight*:  $w_i \leftarrow w_i \cdot \exp[\alpha_m \mathbf{1}(y_i \neq G_m(x_i))]$
- ③ Return *voted classifier*:  $G(x) = \text{sign} \left[ \sum_{m=1}^M \alpha_m G_m(x) \right]$ .

# AdaBoost with Decision Stumps

- After 1 round:

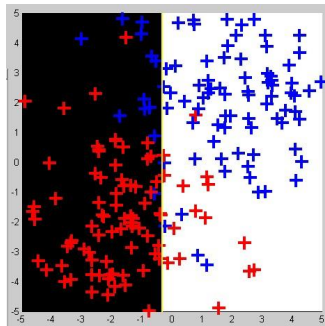


Figure: Plus size represents weight. Blackness represents score for red class.

# AdaBoost with Decision Stumps

- After 3 rounds:




Figure: Plus size represents weight. Blackness represents score for red class.

# AdaBoost with Decision Stumps

- After 120 rounds:

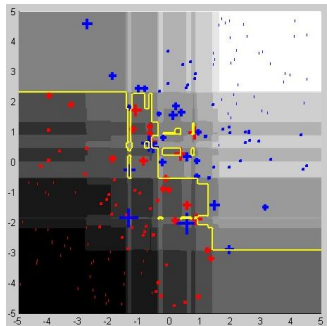
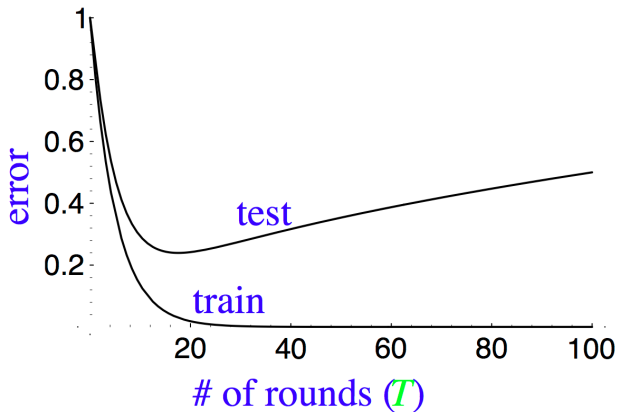


Figure: Plus size represents weight. Blackness represents score for red class.

# Typical Train / Test Learning Curves

- Might expect too many rounds of boosting to overfit:

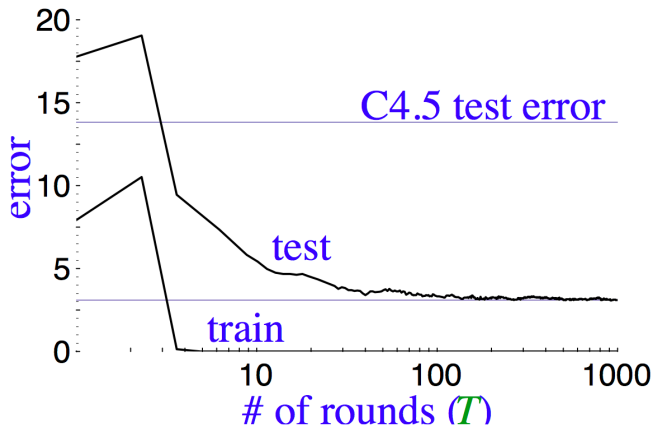


From Rob Schapire's NIPS 2007 Boosting tutorial.



# Learning Curves for AdaBoost

- In typical performance, AdaBoost is surprisingly resistant to overfitting.
- Test continues to improve even after training error is zero!



# Summary

- Shallow decision tree + boosting
  - “best off-the-shelf classifier in the world”—Leo Brieman
  - Used in the first successful real-time face detector (Viola and Jones, 2001)
  - XGBoost: very popular in competitions
- Next week
  - What is the objective function of Adaboost?
  - Generalize to other loss functions.